

# Package ‘ROOT’

November 4, 2025

**Title** Identifying Underrepresented Subpopulations With Interpretable Trees

**Version** 0.0.0.9000

**Description** ROOT (Rashomon set of Optimal Trees) is a framework for learning interpretable binary weight functions represented as sparse decision trees. It constructs a Rashomon set of near-optimal trees and extracts a characteristic tree to summarize patterns. Given trial and target data, the package identifies trial subpopulations that contribute disproportionately to the variance of the target treatment-effect estimate (underrepresented groups).

**License** MIT + file LICENSE

**Encoding** UTF-8

**Roxxygen** list(markdown = TRUE)

**RoxxygenNote** 7.3.3

**Suggests** gbm,  
mlbench,  
testthat (>= 3.0.0),  
knitr,  
rmarkdown,  
ragg

**Config/testthat.edition** 3

**Imports** MASS,

rpart,  
stats,  
withr,  
rpart.plot

**VignetteBuilder** knitr

## Contents

characterize_tree	2
characterizing_underrep	3
choose_feature	4
estimate	5
estimate_dml	6
estimate_dml_single	7
estimate_single	7
gen_S	8

gen_T . . . . .	8
gen_XY . . . . .	9
get_data . . . . .	10
loss_from_objective . . . . .	11
midpoint . . . . .	11
objective_default . . . . .	12
objective_if . . . . .	12
reduce_weight . . . . .	13
ROOT . . . . .	13
split_node . . . . .	15
stratified_kfold . . . . .	16
summary.characterizing_underrep . . . . .	16
summary.ROOT . . . . .	17
train . . . . .	18
train_single . . . . .	19

**Index****20**


---

characterize\_tree      *Fit a shallow decision tree to characterize learned weights w*

---

**Description**

Trains a classification tree on the covariates  $X$  to predict the binary membership  $w$ . This provides an interpretable summary of how the weighted subgroup can be distinguished by  $X$ .

**Usage**

```
characterize_tree(X, w, max_depth = 3)
```

**Arguments**

X	A data frame of covariates (features).
w	A binary vector (0/1 or a factor with two levels) indicating class membership for each observation (e.g., whether an observation is in the selected subgroup).
max_depth	Integer, the maximum tree depth (default 3).

**Details**

The tree is grown using the Gini index (classification) and is not pruned (complexity parameter  $cp = 0$ ), relying solely on  $max\_depth$  to control complexity. This mirrors the default behavior of scikit-learn's `DecisionTreeClassifier(max_depth=...)`. If  $w$  is not already a factor, it will be converted internally. The tree's rules can be interpreted to understand which covariates (and what splits) best separate the two classes defined by  $w$ .

**Value**

An `rpart` object representing the fitted decision tree.

---

**characterizing\_underrep**

*Characterize under-represented subgroups (wraps ROOT)*

---

## Description

Combines an RCT (S=1) and a target dataset (S=0), calls ROOT() to learn a binary selector  $w(x)$ , and (optionally) renders an annotated tree that highlights represented ( $w=1$ ) vs underrepresented ( $w=0$ ) leaves.

## Usage

```
characterizing_underrep(
  DataRCT,
  covariateColName_RCT,
  trtColName_RCT,
  outcomeColName_RCT,
  DataTarget,
  covariateColName_TargetData,
  leaf_proba = 0.25,
  seed = 123,
  num_trees = 10,
  vote_threshold = 2/3,
  explore_proba = 0.05,
  feature_est = "Ridge",
  feature_est_args = list(),
  top_k_trees = FALSE,
  k = 10,
  cutoff = "baseline",
  verbose = FALSE,
  global_objective_fn = objective_default,
  root_plot_tree = TRUE,
  root_plot_args = list(type = 2, extra = 109, under = TRUE, faclen = 0, tweak = 1.1,
    fallen.leaves = TRUE, box.palette = c("pink", "palegreen3"), shadow.col = c("gray"),
    branch.lty = 3, main = "Final Characterized Tree from Rashomon Set"),
  plot_underrep = TRUE,
  keep_threshold = 0.5,
  1X_threshold = NULL,
  plot_main = "Underrepresented Population Characterization Tree"
)
```

## Arguments

DataRCT	data.frame with trial data; must include trtColName_RCT, outcomeColName_RCT, and the covariates named in covariateColName_RCT.
covariateColName_RCT	character vector of covariate column names in DataRCT.
trtColName_RCT	single string: treatment column name in DataRCT (0/1).
outcomeColName_RCT	single string: outcome column name in DataRCT.

**DataTarget** data.frame with target-population covariates (no treatment/outcome required).

**covariateColName\_TargetData** character vector of covariate column names in DataTarget.

**leaf\_proba, seed, num\_trees, vote\_threshold, explore\_proba, feature\_est, feature\_est\_args, top\_k\_trees, k, cutoff, verbose** Passed to ROOT().

**global\_objective\_fn** Global objective/loss function for ROOT (default objective\_default).

**root\_plot\_tree** Logical; pass-through to ROOT(plot\_tree=...).

**root\_plot\_args** Optional list passed to ROOT(plot\_tree\_args=...).

**plot\_underrep** Logical; if TRUE, draws an annotated represented/underrepresented tree.

**keep\_threshold, lX\_threshold** Kept for API compatibility (unused).

**plot\_main** Title for the annotated plot.

### Value

A characterizing\_underrep object with

**root** the fitted ROOT object

**combined** stacked RCT+Target data used for fitting

**leaf\_summary** data.frame with terminal rules and sizes (if a summary tree exists)

**tree\_plot\_root** recorded plot of the ROOT summary tree (if produced)

**tree\_plot\_underrep** recorded plot of the annotated underrep tree (if produced)

choose_feature	<i>Randomly choose a split feature based on provided probabilities</i>
----------------	--

### Description

Given a probability distribution over features (and possibly a "leaf" option), selects one feature at random according to those probabilities.

### Usage

```
choose_feature(split_feature, depth)
```

### Arguments

**split\_feature** A named numeric vector of feature selection probabilities. Names should correspond to feature IDs (and may include a special "leaf" entry).

**depth** Current tree depth (an integer, used for parity with Python implementation but not affecting probabilities in this implementation).

### Value

A single feature name (or "leaf") chosen randomly according to the provided probability weights.

**Note**

The factor  $2^{(0*depth/4)}$  present in the code is effectively 1 (no effect on the first element's weight) and is included only for parity with an equivalent Python implementation. All probabilities are normalized to sum to 1 before sampling.

---

estimate	<i>Compute pseudo-outcome components (a, b) and their product (v)</i>
----------	---

---

**Description**

Using the outputs of the nuisance models, computes intermediate values for the treatment effect estimation via inverse probability weighting (IPW) for the Average Treatment Effect (ATE) in trial sample.

**Usage**

```
estimate(testing_data, outcome, treatment, sample, pi, pi_m, e_m)
```

**Arguments**

testing_data	A data frame of test data (or evaluation data) containing at least the columns for outcome, treatment, and sample indicators.
outcome	Name of the outcome column in testing_data.
treatment	Name of the treatment column in testing_data (0/1).
sample	Name of the sample indicator column in testing_data (0/1).
pi	Numeric scalar, the estimated $P(S = 1)$ (prevalence) from the training data.
pi_m	A fitted model (e.g., gbm) for $P(S = 1   X)$ ; typically from train().
e_m	A fitted model (gbm) for $P(Tr = 1   X, S = 1)$ ; typically from train().

**Details**

Specifically, it computes:

- a: IPW-adjusted outcome difference,  $a_i = S_i \left( \frac{Tr_i Y_i}{p_{t1|x,i}} - \frac{(1-Tr_i)Y_i}{1-p_{t1|x,i}} \right)$ .
- b: Overlap weight factor,  $b_i = \frac{1}{\ell(X_i)}$ , where  $\ell(X) = \frac{P(S=1|X)/\pi}{P(S=0|X)/(1-\pi)}$ .
- v: The pseudo-outcome, defined as  $v_i = a_i \times b_i$ .

The predicted probabilities from pi\_m and e\_m are constrained to [1e-8, 1-1e-8] to avoid instability (extremely small or large probabilities are clamped). If the provided pi is 0 or 1 (indicating no variation in sample inclusion in training), the computation is undefined and an error will be thrown. Ensure that pi\_m and e\_m correspond to models trained on compatible data (same covariates) for accurate predictions.

**Value**

A list with numeric vectors:

v	Pseudo-outcome values for each observation (numeric vector length = nrow(testing_data)).
a	Intermediate "IPW-adjusted outcome" values (same length as v).
b	Overlap weight factors (same length as v).

## See Also

[train](#) for obtaining `pi`, `pi_m`, and `e_m`; [estimate\\_dml](#) for cross-fitted estimation.

`estimate_dml`

*Cross-fitted estimation of pseudo-outcomes (Double ML)*

## Description

Trains nuisance models on each training fold and computes pseudo-outcomes on the corresponding test fold, then aggregates results. Returns only what is needed downstream: the pseudo-outcome table and the aligned evaluation data.

## Usage

```
estimate_dml(data, outcome, treatment, sample, crossfit = 5)
```

## Arguments

<code>data</code>	A data frame containing at least the outcome, treatment, and sample indicator columns.
<code>outcome</code>	Name of the outcome column.
<code>treatment</code>	Name of the treatment column (0/1).
<code>sample</code>	Name of the sample indicator column (0/1).
<code>crossfit</code>	Integer number of folds for cross-fitting ( $\geq 2$ ).

## Value

A list with:

<code>df_v</code>	Data frame with one row per kept observation (indexed by <code>primary_index</code> ), containing: <code>te</code> (pseudo-outcome v), <code>a</code> , <code>b</code> , and squared deviations <code>te_sq</code> , <code>a_sq</code> . Only $S=1$ rows with finite values are kept.
<code>data2</code>	Subset of original data corresponding to <code>df_v\$primary_index</code> .

## Note

Rows with infinite or undefined weights (e.g., where the predicted propensity scores were 0 or 1) are removed from `df_v` (and the corresponding rows in `data2`). The `primary_index` in `df_v` corresponds to the row index in the original data. Squared deviation columns (`te_sq`, `a_sq`) are centered around the mean of `te` and `a` for the  $S=1$  group.

---

<code>estimate_dml_single</code>	<i>Cross-fitted Double ML (single-sample mode)</i>
----------------------------------	--

---

## Description

Runs K-fold cross-fitting to produce pseudo-outcomes for ATE estimation when no sample-membership indicator is available (or has no variation). On each fold, a treatment propensity model is trained on the training split and used to compute pseudo-outcomes on the test split. Results are combined and centered to create variance proxies.

## Usage

```
estimate_dml_single(data, outcome, treatment, crossfit = 5)
```

## Arguments

<code>data</code>	A data frame containing <code>outcome</code> , <code>treatment</code> , and covariates.
<code>outcome</code>	Name of the outcome column.
<code>treatment</code>	Name of the binary treatment indicator column (0/1).
<code>crossfit</code>	Integer number of folds for cross-fitting (default 5; must be $\geq 2$ ).

## Value

A list with:

<code>df_v</code>	Data frame with one row per kept observation, containing: <code>te</code> (pseudo-outcome $v$ ), <code>a</code> , <code>b</code> (all ones), and centered squares <code>te_sq</code> , <code>a_sq</code> , plus <code>primary_index</code> mapping back to the original data rows.
<code>data2</code>	Subset of <code>data</code> corresponding to <code>df_v\$primary_index</code> (i.e., rows kept after cross-fitting and finite checks).

---

<code>estimate_single</code>	<i>Compute single-sample pseudo-outcomes</i>
------------------------------	--

---

## Description

Computes the single-sample pseudo-outcome components for ATE-style estimation:  $a_i = T_i Y_i / e_i - (1 - T_i) Y_i / (1 - e_i)$ , with  $v_i = a_i$  and  $b_i \equiv 1$ . The treatment propensity  $e_i$  is predicted from a supplied model.

## Usage

```
estimate_single(testing_data, outcome, treatment, e_m)
```

## Arguments

<code>testing_data</code>	A data frame containing at least <code>outcome</code> , <code>treatment</code> , and covariates (the latter are used for prediction).
<code>outcome</code>	Name of the outcome column (character).
<code>treatment</code>	Name of the binary treatment indicator column (0/1).
<code>e_m</code>	A fitted <code>glm(binomial)</code> model for $P(T = 1   X)$ ; typically the result of <a href="#">train_single</a> .

**Value**

A list with numeric vectors of length `nrow(testing_data)`:

- `v` Pseudo-outcome values (equal to `a` in single-sample mode).
- `a` IPW-adjusted outcome contrast.
- `b` Vector of ones (no sample-overlap weighting in single-sample mode).

`gen_S`

*Generate sample indicator  $S \sim \text{Bernoulli}(\text{plogis}(a))$*

**Description**

Generates a binary sample inclusion indicator  $S$  for each observation, using a logistic model influenced by a rectangular region in the first two covariates ( $X_0$  and  $X_1$ ).

**Usage**

```
gen_S(X, seed = NULL)
```

**Arguments**

- `X` A data frame of covariates (must contain at least columns  $X_0$  and  $X_1$ ).
- `seed` Optional numeric seed for RNG. If provided, `set.seed(seed + 1)` is invoked for reproducibility. If `NULL` (default), no specific seed is set.

**Details**

The inclusion probability is defined as  $p = \text{plogis}(a)$ , where  $a = 0.25 - 2 * I\{X_0, X_1 \text{ in region } (0.5, 1)\}$ . In other words, observations for which both  $X_0$  and  $X_1$  lie in  $(0.5, 1)$  have a lower odds of being included (due to a negative contribution in the linear predictor). This mirrors a scenario where a specific region in feature space is under-sampled. If a seed is set, it uses `seed + 1` to differentiate from other generators.

**Value**

A data frame with a single column  $S$  of 0/1 values indicating inclusion (1) or exclusion (0).

`gen_T`

*Generate treatment indicator  $Tr \sim \text{Bernoulli}(pi)$*

**Description**

Assigns a treatment indicator for each observation, combining an experimental design for included samples ( $S==1$ ) and an observational assignment for excluded samples ( $S==0$ ).

**Usage**

```
gen_T(X, S, seed = NULL)
```

## Arguments

X	A data frame of covariates.
S	A data frame with column S (0/1 indicating sample inclusion for each observation).
seed	Optional numeric seed for RNG. If provided, set.seed(seed - 1) is used. Default NULL means no explicit seeding.

## Details

For observations with  $S==1$  (in sample), treatment is assigned with probability 0.5 (mimicking a randomized experiment). For those with  $S==0$  (out of sample), treatment probability is  $\text{plogis}(X_0)$ , i.e., it increases with the value of covariate  $X_0$ . The overall assignment probability for each observation is  $\pi_i = S_i * 0.5 + (1 - S_i) * \text{plogis}(X_{0i})$ . If a seed is provided, an offset seed - 1 is used to differentiate from other generation steps.

## Value

A list with two elements:

Tr	A data frame with a single column Tr (treatment assignments 0/1 for each observation).
pi	A numeric vector of length equal to number of observations, giving the treatment probability used for each observation.

gen\_XY

*Generate covariates X and potential outcomes (Y0, Y1)*

## Description

Simulates a regression problem (Friedman #1) and defines a treatment effect. Uses `mlbench.friedman1` to generate X features and a baseline outcome Y0. The treatment potential outcome Y1 is defined as  $Y1 = Y0 + \log(Y0 + 1)$ , introducing a heterogeneous treatment effect.

## Usage

```
gen_XY(n = 1000, seed = NULL)
```

## Arguments

n	Integer or numeric. Number of observations to simulate (must be positive).
seed	Optional. Single numeric value for RNG seed. If provided, a global seed is set for reproducibility. If NULL (default), no seed is set (results will vary on each run).

## Details

The `mlbench.friedman1` function from the **mlbench** package is used to generate 10 independent continuous features and a baseline outcome Y0 with additive noise. The treatment outcome Y1 is defined by adding a non-linear term  $\log(Y0 + 1)$  to the baseline. If a seed is specified, the random number generator state is reset at the start of the function (which affects other random operations).

**Value**

A list with two components:

- X            A data frame of simulated covariates with columns  $X_0, X_1, \dots$  up to  $X_{(p-1)}$ .
- Y            A data frame of potential outcomes with columns  $Y_0$  (baseline outcome) and  $Y_1$  (outcome under treatment).

get\_data

*Convenience wrapper to generate a full simulated dataset*

**Description**

Generates covariates, sample inclusion, treatment assignments, and observed outcomes for a specified sample size. This wraps gen\_XY(), gen\_S(), and gen\_T() in sequence.

**Usage**

```
get_data(n = 1000, seed = NULL)
```

**Arguments**

- n            Integer or numeric. Sample size (number of observations to generate).
- seed        Optional base seed for reproducibility. If provided, internal generators use offsets of this seed to ensure independent randomness. Default NULL means no explicit seeding.

**Details**

This function first generates covariates and potential outcomes with gen\_XY. It then generates S (sample inclusion) and Tr (treatment assignment). The observed outcome  $Y_{obs}$  is computed as  $Y_{obs} = Tr * Y_1 + (1 - Tr) * Y_0$  for each observation.

**Value**

A list with two components:

- data        A data frame of length n containing covariates  $X_0, \dots$ , sample indicator S, treatment indicator Tr, and observed outcome  $Y_{obs}$ .
- Y            A data frame of length n containing the potential outcomes  $Y_0$  and  $Y_1$  for each observation.

**Examples**

```
sim <- get_data(n = 100, seed = 599)
dim(sim$data)    # should be 100 x (p + 3) columns (p features + S + Tr + Yobs)
head(sim$data$Yobs) # observed outcomes
head(sim$Y)      # potential outcomes corresponding to those observations
```

---

`loss_from_objective`     *Backward/fast-path micro-evaluator adaptor*

---

## Description

Wrap a global objective `global_objective_fn(D)` into a splitter-compatible loss function `loss_fn(val, indices, D)` by evaluating `objective_if` on a temporary copy of `D`.

## Usage

```
loss_from_objective(global_objective_fn)
```

## Arguments

`global_objective_fn`

Function of one argument `D` returning a numeric scalar to be minimized (e.g., `objective_default`).

## Value

A function `loss_fn(val, indices, D)` suitable for use in `ROOT` and `split_node`. It sets `w = val` on `indices` (non-mutating), then returns `global_objective_fn(D)`.

---

---

`midpoint`                  *Compute the midpoint of a numeric vector*

---

## Description

Calculates the midpoint defined as  $(\max(X) + \min(X))/2$ , ignoring any NA values.

## Usage

```
midpoint(X)
```

## Arguments

`X`                  A numeric vector.

## Value

A numeric scalar giving the midpoint of the finite values in `X`. If `X` is empty or has no finite values, NA is returned.

<code>objective_default</code>	<i>Default objective: SE proxy of (W)TATE/PATE</i>
--------------------------------	--

### Description

Computes  $\sqrt{\sum_i vsq_i * w_i / (\sum_i w_i)^2}$ . Requires columns vsq and w in D. Minimize this. Supply your own function(D) -> scalar to use a different objective.

### Usage

```
objective_default(D)
```

### Arguments

D data.frame with at least numeric columns vsq and w.

### Value

numeric scalar objective value; Inf if undefined.

<code>objective_if</code>	<i>Helper: evaluate objective after a hypothetical local change</i>
---------------------------	---

### Description

Helper: evaluate objective after a hypothetical local change

### Usage

```
objective_if(val, indices, D, global_objective_fn)
```

### Arguments

val	0/1 assignment to apply
indices	integer or rownames to receive val
D	data.frame used by global_objective_fn
global_objective_fn	function(D)->scalar

### Value

numeric scalar objective after the hypothetical change

---

reduce_weight	<i>Reduce a feature's selection weight by half and renormalize</i>
---------------	--

---

**Description**

Lowers the probability weight of a given feature by 50%, and then re-normalizes the entire probability vector.

**Usage**

```
reduce_weight(fj, split_feature)
```

**Arguments**

- |                            |  |
|----------------------------|--|
| <code>fj</code>            | A feature name (character string) present in the names of <code>split_feature</code> . |
| <code>split_feature</code> | A named numeric vector of probabilities for features (as used in splitting).           |

**Details**

This is typically used when a particular feature split was rejected; the feature's probability is halved to reduce its chance of being chosen again immediately, encouraging exploration of other features. If `fj` is "leaf", its weight is also halved similarly.

**Value**

A numeric vector of the same length as `split_feature`, giving the updated probabilities that sum to 1.

---

ROOT	<i>Ensemble of weighted trees (loss/objective-agnostic) and Rashomon selection</i>
------	--

---

**Description**

Builds multiple weighted trees, then identifies a "Rashomon set" of top-performing trees and aggregates their weight assignments by majority vote.

**Usage**

```
ROOT(
  data,
  outcome,
  treatment,
  sample,
  leaf_proba = 0.25,
  seed = NULL,
  num_trees = 10,
  vote_threshold = 2/3,
  explore_proba = 0.05,
  feature_est = "Ridge",
```

```

feature_est_args = list(),
top_k_trees = FALSE,
k = 10,
cutoff = "baseline",
verbose = FALSE,
global_objective_fn = objective_default,
plot_tree = TRUE,
plot_tree_args = list(type = 2, extra = 109, under = TRUE, faclen = 0, tweak = 1.1,
fallen.leaves = TRUE, box.palette = c("pink", "palegreen3"), shadow.col = c("gray"),
branch.lty = 3, main = "Final Characterized Tree from Rashomon Set")
)

```

## Arguments

<code>data</code>	A data frame containing the dataset (must include outcome, treatment, sample indicator).
<code>outcome</code>	Name of the outcome column in <code>data</code> .
<code>treatment</code>	Name of the treatment indicator column (0/1) in <code>data</code> .
<code>sample</code>	Name of the sample indicator column (0/1) in <code>data</code> . Use <code>NULL</code> for single-sample SATE mode.
<code>leaf_proba</code>	Probability mass for the "leaf" option in each tree (default 0.25).
<code>seed</code>	Integer seed for reproducibility (default <code>NULL</code> ).
<code>num_trees</code>	Number of trees to grow in the forest (default 10).
<code>vote_threshold</code>	Majority vote threshold in (0.5, 1] for final weight=1 (default 2/3).
<code>explore_proba</code>	Probability of exploration at leaves in each tree (default 0.05).
<code>feature_est</code>	"Ridge", "GBM", or a function(X, y, ...) returning a named, non-negative vector of importances; normalized to probabilities (default "Ridge").
<code>feature_est_args</code>	Named list of extra args for a user-supplied <code>feature_est</code> function.
<code>top_k_trees</code>	If TRUE, select top-k trees by objective; else use <code>cutoff</code> (default FALSE).
<code>k</code>	Number of top trees if <code>top_k_trees</code> = TRUE (default 10).
<code>cutoff</code>	If <code>top_k_trees</code> = FALSE, numeric cutoff or "baseline" (default "baseline"). With "baseline", the cutoff is computed by evaluating <code>global_objective_fn</code> on the state with all w=1.
<code>verbose</code>	If TRUE, prints 2 lines with (unweighted and weighted) estimate + SE. Default FALSE.
<code>global_objective_fn</code>	Function <code>function(D) -&gt; numeric</code> scoring the entire state (minimize). Note: Weighted SEs (WATE/WTATE) are printed only when weights are binary (subset-mean). For non-binary weights (custom objectives), SE is omitted by default.
<code>plot_tree</code>	If TRUE, plots the characterized tree (default TRUE). Guarded by <code>interactive()</code> .
<code>plot_tree_args</code>	Named list forwarded to <code>rpart.plot::rpart.plot()</code> .

## Value

S3 object of class "ROOT" with components: `D_rash`, `D_forest`, `w_forest`, `rashomon_set`, `f`, `testing_data`, `tree_plot`, `estimate`

---

<code>split_node</code>	<i>Recursive split builder for weighted tree (internal function)</i>
-------------------------	--

---

## Description

Recursively builds a weighted decision tree to optimize a global objective, using an exploration/exploitation trade-off. Internal; used by ROOT().

## Usage

```
split_node(
  split_feature,
  X,
  D,
  parent_loss,
  depth,
  explore_proba = 0.05,
  choose_feature_fn = choose_feature,
  reduce_weight_fn = reduce_weight,
  global_objective_fn = objective_default,
  max_depth = 8,
  min_leaf_n = 5,
  log_fn = function(...) {
},
  max_rejects_per_node = 1000
)
```

## Arguments

<code>split_feature</code>	Named numeric vector of feature selection probabilities (must include "leaf").
<code>X</code>	Data frame of current observations (includes candidate split feature columns; may include a working copy of weights w).
<code>D</code>	Data frame representing the global state (must include columns w and vsq; row names align to observations).
<code>parent_loss</code>	Numeric, the loss value of the parent node (used to decide if a split improves the objective).
<code>depth</code>	Integer, current tree depth.
<code>explore_proba</code>	Numeric, the probability (between 0 and 1) of flipping the exploit choice at a leaf.
<code>choose_feature_fn</code>	Function to choose next feature (default choose_feature).
<code>reduce_weight_fn</code>	Function to penalize last-tried feature on rejected split (default reduce_weight).
<code>global_objective_fn</code>	Function function(D) -> numeric scoring the <b>entire</b> state.
<code>max_depth</code>	Integer max depth (stop and make leaf at this depth).
<code>min_leaf_n</code>	Integer min rows to attempt a split; else make leaf.
<code>log_fn</code>	Function for logging; default no-op.
<code>max_rejects_per_node</code>	Safety budget of rejected splits before forcing a leaf.

**Value**

A list representing the (sub)tree; includes updated D and local objective.

stratified_kfold	<i>Stratified K-fold index generator</i>
------------------	--

**Description**

Splits indices into K folds while preserving the class distribution of a binary factor. This mimics scikit-learn's `StratifiedKFold`, ensuring each fold has a representative ratio of the two classes in S.

**Usage**

```
stratified_kfold(S, K = 5)
```

**Arguments**

- |   |  |
|---|--|
| S | A vector or factor indicating class membership (typically 0/1 or two-class factor) for stratification.                     |
| K | Integer number of folds ( $K \geq 2$ ). If K is larger than the number of observations, it will be reduced to that number. |

**Details**

The function deterministically allocates indices to folds by class. For each class in S, indices are cyclically assigned to folds to balance counts. If K == 1, a single fold containing all indices is returned (though typically K should be  $\geq 2$  for cross-validation).

**Value**

A list of length K, where each element is an integer vector of row indices assigned to that fold. The union of all folds equals 1:length(S), and folds are roughly equal in size.

summary.characterizing_underrep	<i>Summarize a characterizing_underrep fit</i>
---------------------------------	--

**Description**

Prints the ROOT summary (un/weighted estimates with standard errors; the *weighted* SE is omitted when a custom `global_objective_fn` was used in `ROOT()`) and a brief overview of terminal rules from the annotated summary tree, if available.

**Usage**

```
## S3 method for class 'characterizing_underrep'
summary(object, ...)
```

**Arguments**

- |        |  |
|--------|--|
| object | A characterizing_underrep object.      |
| ...    | Unused; included for S3 compatibility. |

**Details**

Delegates core statistics to `summary(object$root)`; previews up to ten terminal rules when a summary tree exists, and reports plot availability.

**Value**

`object`, invisibly.

summary.ROOT	<i>Summarize a ROOT fit</i>
--------------	-----------------------------

**Description**

Summarizes a ROOT object by reporting the primary estimands and key model diagnostics. The first lines report:

1. the **unweighted** estimate (ATE in RCT for single-sample or TATE for two-sample) and its **standard error (SE)**;
2. the **weighted** estimate (WATE in RCT or WTATE using `w_opt`) and its **SE** whenever `w_opt` is effectively *binary* (subset-mean SE); if `w_opt` is non-binary, the SE is omitted with a note.

Subsequent lines describe the estimand type, number of trees, size of the Rashomon set, presence of a summary tree, covariate count, observation count, baseline loss, selected-tree losses, and the proportion kept by `w_opt`.

**Usage**

```
## S3 method for class 'ROOT'
summary(object, ...)
```

**Arguments**

- |        |   |
|--------|---|
| object | A ROOT object returned by <code>ROOT()</code> . |
| ...    | Unused; included for S3 compatibility.          |

**Details**

This method prefers pre-computed estimates. If unavailable, it recomputes:

- Unweighted effect as  $\bar{v}$  over the analysis set (all rows in single-sample;  $S = 1$  in two-sample);
- Unweighted SE as  $\sqrt{\frac{1}{n(n-1)} \sum (v_i - \bar{v})^2} \sqrt{\sum((v - v_{\bar{A}})^2) / (n * (n - 1))}$ ;
- Weighted effect when `w_opt` is binary, with  $A = \{i : w_i = 1\}$ , i.e.  $\bar{v}_A = \frac{1}{n_A} \sum_{i \in A} v_i$ ;
- Weighted SE (WTATE/WATE), for binary `w_opt`, as  $\sqrt{\frac{1}{n_A(n_A-1)} \sum_{i \in A} (v_i - \bar{v}_A)^2} \sqrt{\sum((v[w == 1] - v_{\bar{A}})^2) / (n_A * (n_A - 1))}$ .

**Value**

The input object, invisibly. Printed output is a human-readable summary.

**train***Train nuisance models for weighting***Description**

Fits models to estimate sampling and treatment propensities on training data. Specifically, it computes:

- *pi*: The prevalence of sample inclusion (estimated as mean of *S*).
- *pi\_m*: A logistic regression model for  $P(S = 1 | X)$  using all covariates.
- *e\_m*: A logistic regression model for  $P(Tr = 1 | X, S = 1)$ , fit only on the subset where  $S == 1$ .

**Usage**

```
train(training_data, outcome, treatment, sample)
```

**Arguments**

<i>training_data</i>	A data frame containing the training dataset.
<i>outcome</i>	Name of the outcome column (typically observed outcome, e.g. "Yobs").
<i>treatment</i>	Name of the treatment indicator column (e.g. "Tr").
<i>sample</i>	Name of the sample inclusion indicator column (e.g. "S").

**Details**

This function uses simple logistic regression (`glm` with logit link) to estimate the necessary nuisance parameters for weighting. It requires that both *S* and *Tr* have variation (both 0 and 1 must be present in training data); if not, the fitting is not possible and an error is raised. All covariates other than the specified outcome, treatment, and sample columns are used as predictors.

**Value**

A list with components:

<i>pi</i>	Numeric scalar giving the overall sample inclusion rate $P(S = 1)$ in the training data.
<i>pi_m</i>	A fitted <code>glm</code> model (binomial family) for $P(S = 1   X)$ .
<i>e_m</i>	A fitted <code>glm</code> model (binomial family) for $P(Tr = 1   X, S = 1)$ .

---

train_single	<i>Train treatment propensity model (single-sample mode)</i>
--------------	--

---

## Description

Fits a logistic regression for  $P(T = 1 | X)$  on the provided training data. Used by the single-sample Double ML path where no sample-selection model is needed.

## Usage

```
train_single(training_data, outcome, treatment)
```

## Arguments

- training\_data** A data frame containing the outcome, treatment, and covariates. Only **treatment** and covariates are used for fitting.
- outcome** Name of the outcome column (character). Present for a consistent signature but not used in this function.
- treatment** Name of the binary treatment indicator column (0/1).

## Value

A list with one element:

- e\_m** A `glm` (binomial) object for the treatment propensity model.

# Index

characterize\_tree, 2  
characterizing\_underrep, 3  
choose\_feature, 4  
  
estimate, 5  
estimate\_dml, 6, 6  
estimate\_dml\_single, 7  
estimate\_single, 7  
  
gen\_S, 8  
gen\_T, 8  
gen\_XY, 9  
get\_data, 10  
  
loss\_from\_objective, 11  
  
midpoint, 11  
  
objective\_default, 11, 12  
objective\_if, 11, 12  
  
reduce\_weight, 13  
ROOT, 11, 13  
  
split\_node, 11, 15  
stratified\_kfold, 16  
summary.characterizing\_underrep, 16  
summary.ROOT, 17  
  
train, 6, 18  
train\_single, 7, 19