

# Package ‘ROOT’

November 4, 2025

**Title** Identifying Underrepresented Subpopulations With Interpretable Trees

**Version** 0.0.0.9000

**Description** ROOT (Rashomon set of Optimal Trees) is a framework for learning interpretable binary weight functions represented as sparse decision trees. It constructs a Rashomon set of near-optimal trees and extracts a characteristic tree to summarize patterns. Given trial and target data, the package identifies trial subpopulations that contribute disproportionately to the variance of the target treatment-effect estimate (underrepresented groups).

**License** MIT + file LICENSE

**Encoding** UTF-8

**Roxxygen** list(markdown = TRUE)

**RoxxygenNote** 7.3.3

**Suggests** gbm,  
mlbench,  
testthat (>= 3.0.0),  
knitr,  
rmarkdown,  
ragg

**Config/testthat.edition** 3

**Imports** MASS,

rpart,  
stats,  
withr,  
rpart.plot

**VignetteBuilder** knitr

## Contents

characterize_tree	2
characterizing_underrep	3
choose_feature	4
estimate	5
estimate_dml	6
estimate_dml_single	7
estimate_single	7
gen_S	8

gen_T . . . . .	8
gen_XY . . . . .	9
get_data . . . . .	10
loss . . . . .	11
loss_from_objective . . . . .	11
midpoint . . . . .	12
objective_default . . . . .	12
objective_if . . . . .	13
reduce_weight . . . . .	13
ROOT . . . . .	14
split_node . . . . .	15
stratified_kfold . . . . .	17
summary.characterizing_underrep . . . . .	17
summary.ROOT . . . . .	18
train . . . . .	19
train_single . . . . .	20

**Index****21**


---

**characterize\_tree**      *Fit a shallow decision tree to characterize learned weights w*

---

**Description**

Trains a classification tree on the covariates  $X$  to predict the binary membership  $w$ . This provides an interpretable summary of how the weighted subgroup can be distinguished by  $X$ .

**Usage**

```
characterize_tree(X, w, max_depth = 3)
```

**Arguments**

- X**      A data frame of covariates (features).
- w**      A binary vector (0/1 or a factor with two levels) indicating class membership for each observation (e.g., whether an observation is in the selected subgroup).
- max\_depth**      Integer, the maximum tree depth (default 3).

**Details**

The tree is grown using the Gini index (classification) and is not pruned (complexity parameter  $cp = 0$ ), relying solely on  $\text{max\_depth}$  to control complexity. This mirrors the default behavior of scikit-learn's `DecisionTreeClassifier(max_depth=...)`. If  $w$  is not already a factor, it will be converted internally. The tree's rules can be interpreted to understand which covariates (and what splits) best separate the two classes defined by  $w$ .

**Value**

An `rpart` object representing the fitted decision tree.

---

**characterizing\_underrep**

*Characterize under-represented subgroups (wraps ROOT)*

---

## Description

Combines an RCT (S=1) and a target dataset (S=0), calls ROOT() to learn a binary selector  $w(x)$ , and (optionally) renders an annotated tree that highlights represented ( $w=1$ ) vs underrepresented ( $w=0$ ) leaves.

## Usage

```
characterizing_underrep(
  DataRCT,
  covariate_DataRCT,
  treatment_DataRCT,
  outcome_DataRCT,
  DataTarget,
  covariate_DataTarget,
  leaf_proba = 0.25,
  seed = 123,
  num_trees = 10,
  vote_threshold = 2/3,
  explore_proba = 0.05,
  feature_est = "Ridge",
  feature_est_args = list(),
  top_k_trees = FALSE,
  k = 10,
  cutoff = "baseline",
  verbose = FALSE,
  objective_fn = objective_default,
  loss_fn = NULL,
  root_plot_tree = TRUE,
  root_plot_args = list(type = 2, extra = 109, under = TRUE, faclen = 0, tweak = 1.1,
    fallen.leaves = TRUE, box.palette = c("pink", "palegreen3"), shadow.col = c("gray"),
    branch.lty = 3, main = "Final Characterized Tree from Rashomon Set"),
  plot_underrep = TRUE,
  keep_threshold = 0.5,
  lX_threshold = NULL,
  plot_main = "Underrepresented Population Characterization Tree"
)
```

## Arguments

<b>DataRCT</b>	data.frame with trial data; must include treatment_DataRCT, outcome_DataRCT, and the covariates named in covariate_DataRCT.
<b>covariate_DataRCT</b>	character vector of covariate column names in DataRCT.
<b>treatment_DataRCT</b>	single string: treatment column name in DataRCT (0/1).

```

outcome_DataRCT           single string: outcome column name in DataRCT.
DataTarget                 data.frame with target-population covariates (no treatment/outcome required).
covariate_DataTarget       character vector of covariate column names in DataTarget.
leaf_proba, seed, num_trees, vote_threshold, explore_proba, feature_est,
feature_est_args, top_k_trees, k, cutoff, verbose
                           Passed to ROOT().
objective_fn               Objective function for ROOT (default objective_default).
loss_fn                    Optional loss micro-evaluator; if NULL, ROOT will wrap objective_fn.
root_plot_tree             Logical; pass-through to ROOT(plot_tree=...).
root_plot_args              Optional list passed to ROOT(plot_tree_args=...).
plot_underrep              Logical; if TRUE, draws an annotated represented/underrepresented tree.
keep_threshold, lX_threshold
                           Kept for API compatibility (unused).
plot_main                  Title for the annotated plot.

```

### **Value**

A `characterizing_underrep` object with

```

root                     the fitted ROOT object
combined                 stacked RCT+Target data used for fitting
leaf_summary              data.frame with terminal rules and sizes (if a summary tree exists)
tree_plot_root            recorded plot of the ROOT summary tree (if produced)
tree_plot_underrep
                           recorded plot of the annotated underrep tree (if produced)

```

<code>choose_feature</code>	<i>Randomly choose a split feature based on provided probabilities</i>
-----------------------------	--

### **Description**

Given a probability distribution over features (and possibly a "leaf" option), selects one feature at random according to those probabilities.

### **Usage**

```
choose_feature(split_feature, depth)
```

### **Arguments**

<code>split_feature</code>	A named numeric vector of feature selection probabilities. Names should correspond to feature IDs (and may include a special "leaf" entry).
<code>depth</code>	Current tree depth (an integer, used for parity with Python implementation but not affecting probabilities in this implementation).

**Value**

A single feature name (or "leaf") chosen randomly according to the provided probability weights.

**Note**

The factor  $2^{(0*depth/4)}$  present in the code is effectively 1 (no effect on the first element's weight) and is included only for parity with an equivalent Python implementation. All probabilities are normalized to sum to 1 before sampling.

estimate

*Compute pseudo-outcome components (a, b) and their product (v)*

**Description**

Using the outputs of the nuisance models, computes intermediate values for the treatment effect estimation via inverse probability weighting (IPW) for the Average Treatment Effect (ATE) in trial sample.

**Usage**

```
estimate(testing_data, outcome, treatment, sample, pi, pi_m, e_m)
```

**Arguments**

testing_data	A data frame of test data (or evaluation data) containing at least the columns for outcome, treatment, and sample indicators.
outcome	Name of the outcome column in testing_data.
treatment	Name of the treatment column in testing_data (0/1).
sample	Name of the sample indicator column in testing_data (0/1).
pi	Numeric scalar, the estimated $P(S = 1)$ (prevalence) from the training data.
pi_m	A fitted model (e.g., glm) for $P(S = 1   X)$ ; typically from train().
e_m	A fitted model (glm) for $P(Tr = 1   X, S = 1)$ ; typically from train().

**Details**

Specifically, it computes:

- a: IPW-adjusted outcome difference,  $a_i = S_i \left( \frac{Tr_i Y_i}{p_{t1|x,i}} - \frac{(1-Tr_i)Y_i}{1-p_{t1|x,i}} \right)$ .
- b: Overlap weight factor,  $b_i = \frac{1}{\ell(X_i)}$ , where  $\ell(X) = \frac{P(S=1|X)/\pi}{P(S=0|X)/(1-\pi)}$ .
- v: The pseudo-outcome, defined as  $v_i = a_i \times b_i$ .

The predicted probabilities from pi\_m and e\_m are constrained to [1e-8, 1-1e-8] to avoid instability (extremely small or large probabilities are clamped). If the provided pi is 0 or 1 (indicating no variation in sample inclusion in training), the computation is undefined and an error will be thrown. Ensure that pi\_m and e\_m correspond to models trained on compatible data (same covariates) for accurate predictions.

**Value**

A list with numeric vectors:

- v Pseudo-outcome values for each observation (numeric vector length = nrow(testing\_data)).
- a Intermediate "IPW-adjusted outcome" values (same length as v).
- b Overlap weight factors (same length as v).

**See Also**

[train](#) for obtaining `pi`, `pi_m`, and `e_m`; [estimate\\_dml](#) for cross-fitted estimation.

`estimate_dml`

*Cross-fitted estimation of pseudo-outcomes (Double ML)*

**Description**

Trains nuisance models on each training fold and computes pseudo-outcomes on the corresponding test fold, then aggregates results. Returns only what is needed downstream: the pseudo-outcome table and the aligned evaluation data.

**Usage**

```
estimate_dml(data, outcome, treatment, sample, crossfit = 5)
```

**Arguments**

- data A data frame containing at least the outcome, treatment, and sample indicator columns.
- outcome Name of the outcome column.
- treatment Name of the treatment column (0/1).
- sample Name of the sample indicator column (0/1).
- crossfit Integer number of folds for cross-fitting ( $\geq 2$ ).

**Value**

A list with:

- df\_v Data frame with one row per kept observation (indexed by `primary_index`), containing: te (pseudo-outcome v), a, b, and squared deviations te\_sq, a\_sq. Only S==1 rows with finite values are kept.
- data2 Subset of original data corresponding to `df_v$primary_index`.

**Note**

Rows with infinite or undefined weights (e.g., where the predicted propensity scores were 0 or 1) are removed from `df_v` (and the corresponding rows in `data2`). The `primary_index` in `df_v` corresponds to the row index in the original data. Squared deviation columns (`te_sq`, `a_sq`) are centered around the mean of te and a for the S==1 group.

---

<code>estimate_dml_single</code>	<i>Cross-fitted Double ML (single-sample mode)</i>
----------------------------------	--

---

## Description

Runs K-fold cross-fitting to produce pseudo-outcomes for ATE estimation when no sample-membership indicator is available (or has no variation). On each fold, a treatment propensity model is trained on the training split and used to compute pseudo-outcomes on the test split. Results are combined and centered to create variance proxies.

## Usage

```
estimate_dml_single(data, outcome, treatment, crossfit = 5)
```

## Arguments

<code>data</code>	A data frame containing <code>outcome</code> , <code>treatment</code> , and covariates.
<code>outcome</code>	Name of the outcome column.
<code>treatment</code>	Name of the binary treatment indicator column (0/1).
<code>crossfit</code>	Integer number of folds for cross-fitting (default 5; must be $\geq 2$ ).

## Value

A list with:

<code>df_v</code>	Data frame with one row per kept observation, containing: <code>te</code> (pseudo-outcome $v$ ), <code>a</code> , <code>b</code> (all ones), and centered squares <code>te_sq</code> , <code>a_sq</code> , plus <code>primary_index</code> mapping back to the original data rows.
<code>data2</code>	Subset of <code>data</code> corresponding to <code>df_v\$primary_index</code> (i.e., rows kept after cross-fitting and finite checks).

---

<code>estimate_single</code>	<i>Compute single-sample pseudo-outcomes</i>
------------------------------	--

---

## Description

Computes the single-sample pseudo-outcome components for ATE-style estimation:  $a_i = T_i Y_i / e_i - (1 - T_i) Y_i / (1 - e_i)$ , with  $v_i = a_i$  and  $b_i \equiv 1$ . The treatment propensity  $e_i$  is predicted from a supplied model.

## Usage

```
estimate_single(testing_data, outcome, treatment, e_m)
```

## Arguments

<code>testing_data</code>	A data frame containing at least <code>outcome</code> , <code>treatment</code> , and covariates (the latter are used for prediction).
<code>outcome</code>	Name of the outcome column (character).
<code>treatment</code>	Name of the binary treatment indicator column (0/1).
<code>e_m</code>	A fitted <code>glm(binomial)</code> model for $P(T = 1   X)$ ; typically the result of <a href="#">train_single</a> .

**Value**

A list with numeric vectors of length `nrow(testing_data)`:

- `v` Pseudo-outcome values (equal to `a` in single-sample mode).
- `a` IPW-adjusted outcome contrast.
- `b` Vector of ones (no sample-overlap weighting in single-sample mode).

`gen_S`

*Generate sample indicator  $S \sim \text{Bernoulli}(\text{plogis}(a))$*

**Description**

Generates a binary sample inclusion indicator  $S$  for each observation, using a logistic model influenced by a rectangular region in the first two covariates ( $X_0$  and  $X_1$ ).

**Usage**

```
gen_S(X, seed = NULL)
```

**Arguments**

- `X` A data frame of covariates (must contain at least columns  $X_0$  and  $X_1$ ).
- `seed` Optional numeric seed for RNG. If provided, `set.seed(seed + 1)` is invoked for reproducibility. If `NULL` (default), no specific seed is set.

**Details**

The inclusion probability is defined as  $p = \text{plogis}(a)$ , where  $a = 0.25 - 2 * I\{X_0, X_1 \text{ in region } (0.5, 1)\}$ . In other words, observations for which both  $X_0$  and  $X_1$  lie in  $(0.5, 1)$  have a lower odds of being included (due to a negative contribution in the linear predictor). This mirrors a scenario where a specific region in feature space is under-sampled. If a seed is set, it uses `seed + 1` to differentiate from other generators.

**Value**

A data frame with a single column  $S$  of 0/1 values indicating inclusion (1) or exclusion (0).

`gen_T`

*Generate treatment indicator  $Tr \sim \text{Bernoulli}(pi)$*

**Description**

Assigns a treatment indicator for each observation, combining an experimental design for included samples ( $S==1$ ) and an observational assignment for excluded samples ( $S==0$ ).

**Usage**

```
gen_T(X, S, seed = NULL)
```

## Arguments

X	A data frame of covariates.
S	A data frame with column S (0/1 indicating sample inclusion for each observation).
seed	Optional numeric seed for RNG. If provided, set.seed(seed - 1) is used. Default NULL means no explicit seeding.

## Details

For observations with  $S==1$  (in sample), treatment is assigned with probability 0.5 (mimicking a randomized experiment). For those with  $S==0$  (out of sample), treatment probability is  $\text{plogis}(X_0)$ , i.e., it increases with the value of covariate  $X_0$ . The overall assignment probability for each observation is  $\pi_i = S_i * 0.5 + (1 - S_i) * \text{plogis}(X_{0i})$ . If a seed is provided, an offset seed - 1 is used to differentiate from other generation steps.

## Value

A list with two elements:

Tr	A data frame with a single column Tr (treatment assignments 0/1 for each observation).
pi	A numeric vector of length equal to number of observations, giving the treatment probability used for each observation.

gen\_XY

*Generate covariates X and potential outcomes (Y0, Y1)*

## Description

Simulates a regression problem (Friedman #1) and defines a treatment effect. Uses `mlbench.friedman1` to generate X features and a baseline outcome Y0. The treatment potential outcome Y1 is defined as  $Y1 = Y0 + \log(Y0 + 1)$ , introducing a heterogeneous treatment effect.

## Usage

```
gen_XY(n = 1000, seed = NULL)
```

## Arguments

n	Integer or numeric. Number of observations to simulate (must be positive).
seed	Optional. Single numeric value for RNG seed. If provided, a global seed is set for reproducibility. If NULL (default), no seed is set (results will vary on each run).

## Details

The `mlbench.friedman1` function from the **mlbench** package is used to generate 10 independent continuous features and a baseline outcome Y0 with additive noise. The treatment outcome Y1 is defined by adding a non-linear term  $\log(Y0 + 1)$  to the baseline. If a seed is specified, the random number generator state is reset at the start of the function (which affects other random operations).

**Value**

A list with two components:

- X A data frame of simulated covariates with columns  $X_0, X_1, \dots$  up to  $X_{(p-1)}$ .
- Y A data frame of potential outcomes with columns  $Y_0$  (baseline outcome) and  $Y_1$  (outcome under treatment).

get\_data

*Convenience wrapper to generate a full simulated dataset*

**Description**

Generates covariates, sample inclusion, treatment assignments, and observed outcomes for a specified sample size. This wraps gen\_XY(), gen\_S(), and gen\_T() in sequence.

**Usage**

```
get_data(n = 1000, seed = NULL)
```

**Arguments**

- n Integer or numeric. Sample size (number of observations to generate).
- seed Optional base seed for reproducibility. If provided, internal generators use offsets of this seed to ensure independent randomness. Default NULL means no explicit seeding.

**Details**

This function first generates covariates and potential outcomes with gen\_XY. It then generates S (sample inclusion) and Tr (treatment assignment). The observed outcome  $Y_{obs}$  is computed as  $Y_{obs} = Tr * Y_1 + (1 - Tr) * Y_0$  for each observation.

**Value**

A list with two components:

- data A data frame of length n containing covariates  $X_0, \dots$ , sample indicator S, treatment indicator Tr, and observed outcome  $Y_{obs}$ .
- Y A data frame of length n containing the potential outcomes  $Y_0$  and  $Y_1$  for each observation.

**Examples**

```
sim <- get_data(n = 100, seed = 599)
dim(sim$data)    # should be 100 x (p + 3) columns (p features + S + Tr + Yobs)
head(sim$data$Yobs) # observed outcomes
head(sim$Y)      # potential outcomes corresponding to those observations
```

---

loss	<i>Evaluate splitting objective loss for a given weight assignment</i>
------	--

---

**Description**

Computes the approximate standard error (loss) of the treatment effect estimate under a hypothetical assignment of weights w for specified rows.

**Usage**

```
loss(val, indices, D)
```

**Arguments**

val	Numeric scalar (must be 0 or 1). The weight value to assign (0 = exclude, 1 = include).
indices	Indices or row names in D for which the weight should be set to val. Can be a numeric vector of row positions or a character vector of row names.
D	A data frame containing at least columns vsq (squared pseudo-outcome) and w (current weights).

**Value**

Numeric value representing the loss, defined as  $\sqrt{\sum_i vsq_i * w_i / (\sum_i w_i)^2}$ . Returns Inf if the denominator is 0 or if the result is not a number.

**Note**

This function mimics the behavior of numpy's `nan_to_num(..., nan=Inf)` by returning Inf when the computation is undefined (e.g., no weights selected). It is used internally to decide whether a proposed split improves the objective.

---

loss_from_objective	<i>Backward/fast-path micro-evaluator adaptor</i>
---------------------	---

---

**Description**

Wrap a global objective `objective_fn(D)` into a splitter-compatible loss function `loss_fn(val, indices, D)` by evaluating `objective_if` on a temporary copy of D.

**Usage**

```
loss_from_objective(objective_fn)
```

**Arguments**

objective_fn	Function of one argument D returning a numeric scalar to be minimized (e.g., <code>objective_default</code> ).
--------------	--

**Value**

A function `loss_fn(val, indices, D)` suitable for use in `ROOT` and `split_node`. It sets `w = val` on `indices` (non-mutating), then returns `objective_fn(D)`.

`midpoint`*Compute the midpoint of a numeric vector***Description**

Calculates the midpoint defined as  $(\max(X) + \min(X))/2$ , ignoring any NA values.

**Usage**

```
midpoint(X)
```

**Arguments**

<code>X</code>	A numeric vector.
----------------	-------------------

**Value**

A numeric scalar giving the midpoint of the finite values in `X`. If `X` is empty or has no finite values, `NA` is returned.

`objective_default`*Default objective: SE proxy of (W)TATE/PATE***Description**

Computes  $\sqrt{\sum_i vsq_i * w_i / (\sum_i w_i)^2}$ . Requires columns `vsq` and `w` in `D`. Minimize this. Supply your own `function(D) -> scalar` to use a different objective.

**Usage**

```
objective_default(D)
```

**Arguments**

<code>D</code>	data.frame with at least numeric columns <code>vsq</code> and <code>w</code> .
----------------	--

**Value**

numeric scalar objective value; `Inf` if undefined.

**objective\_if***Helper: evaluate objective after a hypothetical local change***Description**

Helper: evaluate objective after a hypothetical local change

**Usage**

```
objective_if(val, indices, D, objective_fn)
```

**Arguments**

val	0/1 assignment to apply
indices	integer or rownames to receive val
D	data.frame used by objective_fn
objective_fn	function(D)->scalar

**Value**

numeric scalar objective after the hypothetical change

**reduce\_weight***Reduce a feature's selection weight by half and renormalize***Description**

Lowers the probability weight of a given feature by 50%, and then re-normalizes the entire probability vector.

**Usage**

```
reduce_weight(fj, split_feature)
```

**Arguments**

fj	A feature name (character string) present in the names of split_feature.
split_feature	A named numeric vector of probabilities for features (as used in splitting).

**Details**

This is typically used when a particular feature split was rejected; the feature's probability is halved to reduce its chance of being chosen again immediately, encouraging exploration of other features. If fj is "leaf", its weight is also halved similarly.

**Value**

A numeric vector of the same length as split\_feature, giving the updated probabilities that sum to 1.

---

ROOT	<i>Ensemble of weighted trees (loss/objective-agnostic) and Rashomon selection</i>
------	--

---

## Description

Builds multiple weighted trees, then identifies a "Rashomon set" of top-performing trees and aggregates their weight assignments by majority vote.

## Usage

```
ROOT(
  data,
  outcome,
  treatment,
  sample,
  leaf_proba = 0.25,
  seed = NULL,
  num_trees = 10,
  vote_threshold = 2/3,
  explore_proba = 0.05,
  feature_est = "Ridge",
  feature_est_args = list(),
  top_k_trees = FALSE,
  k = 10,
  cutoff = "baseline",
  verbose = FALSE,
  objective_fn = objective_default,
  loss_fn = NULL,
  plot_tree = TRUE,
  plot_tree_args = list(type = 2, extra = 109, under = TRUE, faclen = 0, tweak = 1.1,
    fallen.leaves = TRUE, box.palette = c("pink", "palegreen3"), shadow.col = c("gray"),
    branch.lty = 3, main = "Final Characterized Tree from Rashomon Set")
)
```

## Arguments

<code>data</code>	A data frame containing the dataset (must include outcome, treatment, sample indicator).
<code>outcome</code>	Name of the outcome column in <code>data</code> .
<code>treatment</code>	Name of the treatment indicator column (0/1) in <code>data</code> .
<code>sample</code>	Name of the sample indicator column (0/1) in <code>data</code> . Use <code>NULL</code> for single-sample SATE mode.
<code>leaf_proba</code>	Probability mass for the "leaf" option in each tree (default 0.25).
<code>seed</code>	Integer seed for reproducibility (default <code>NULL</code> ).
<code>num_trees</code>	Number of trees to grow in the forest (default 10).
<code>vote_threshold</code>	Majority vote threshold in (0.5, 1] for final weight=1 (default 2/3).
<code>explore_proba</code>	Probability of exploration at leaves in each tree (default 0.05).

<code>feature_est</code>	"Ridge", "GBM", or a function(X, y, ...) returning a named, non-negative vector of importances; normalized to probabilities (default "Ridge").
<code>feature_est_args</code>	Named list of extra args for a user-supplied <code>feature_est</code> function.
<code>top_k_trees</code>	If TRUE, select top-k trees by objective; else use <code>cutoff</code> (default FALSE).
<code>k</code>	Number of top trees if <code>top_k_trees</code> = TRUE (default 10).
<code>cutoff</code>	If <code>top_k_trees</code> = FALSE, numeric cutoff or "baseline" (default "baseline"). With "baseline", the cutoff is computed by evaluating <code>objective_fn</code> on the state with all w=1.
<code>verbose</code>	If TRUE, prints 2 lines with (unweighted and weighted) estimate + SD. Default FALSE.
<code>objective_fn</code>	Function <code>function(D) -&gt; numeric</code> that scores the <b>entire state</b> (minimize). Default <code>objective_default()</code> reproduces prior behavior (SE proxy of PATE/TATE using <code>vsq</code> and <code>w</code> ).
<code>loss_fn</code>	Optional micro-evaluator function( <code>val, indices, D</code> ) -> numeric that returns the objective after hypothetically setting <code>w=val</code> on <code>indices</code> . If NULL, <code>ROOT</code> wraps <code>objective_fn</code> via <code>loss_from_objective(objective_fn)</code> .
<code>plot_tree</code>	If TRUE, plots the characterized tree (default TRUE). Guarded by <code>interactive()</code> .
<code>plot_tree_args</code>	Named list forwarded to <code>rpart::rpart.plot()</code> .

**Value**

S3 object of class "ROOT" with components: `D_rash`, `D_forest`, `w_forest`, `rashomon_set`, `f`, `testing_data`, `tree_plot`, `estimate`

<code>split_node</code>	<i>Recursive split builder for weighted tree (internal function)</i>	<i>Recursive split builder for weighted tree (internal function)</i>
-------------------------	--	--

**Description**

Recursively builds a weighted decision tree to optimize an objective function, using an exploration/exploitation trade-off. This function is internal and is used by `tree_opt()` / `ROOT()` to construct a single tree.

**Usage**

```
split_node(
  split_feature,
  X,
  D,
  parent_loss,
  depth,
  explore_proba = 0.05,
  choose_feature_fn = choose_feature,
  loss_fn = loss,
  reduce_weight_fn = reduce_weight,
  objective_fn = objective_default,
```

```

max_depth = 8,
min_leaf_n = 5,
log_fn = function(...) {
},
max_rejects_per_node = 1000
)

```

## Arguments

<code>split_feature</code>	Named numeric vector of feature selection probabilities (should include a "leaf" option).
<code>X</code>	Data frame of current observations (must include at least the feature chosen for splitting if applicable; may also include a working copy of weights <code>w</code> ).
<code>D</code>	Data frame representing the global state (must include columns <code>w</code> for weights and <code>vsq</code> for squared pseudo-outcomes, with row names aligning to observations).
<code>parent_loss</code>	Numeric, the loss value of the parent node (used to decide if a split improves the objective).
<code>depth</code>	Integer, current depth of the node in the tree.
<code>explore_proba</code>	Numeric in $[0, 1]$ , probability of exploring the suboptimal assignment at a leaf (randomly flipping the chosen weight).
<code>choose_feature_fn</code>	Function to choose the next feature to split; default <code>choose_feature</code> .
<code>loss_fn</code>	Function to compute loss given a tentative weight assignment; if <code>NULL</code> , it is set via <code>loss_from_objective(objective_fn)</code> .
<code>reduce_weight_fn</code>	Function to adjust feature probabilities when a split is rejected; default <code>reduce_weight</code> .
<code>objective_fn</code>	Function that maps <code>D</code> $\rightarrow$ scalar objective; used for node summaries and to derive <code>loss_fn</code> when <code>loss_fn</code> is <code>NULL</code> .
<code>max_depth</code>	Integer maximum depth of the tree. If the current depth equals <code>max_depth</code> , the node is made a leaf.
<code>min_leaf_n</code>	Integer minimum number of observations required to attempt a split. If <code>X</code> has $<= \text{min\_leaf\_n}$ rows, the node becomes a leaf.
<code>log_fn</code>	Function for logging debug messages. Default no-op.
<code>max_rejects_per_node</code>	Integer. A safety budget for how many times a node may <i>reject</i> non-improving splits before the algorithm gives up and finalizes the node as a leaf. Each rejection halves the probability of the last-tried feature (via <code>reduce_weight_fn</code> ) and samples a new candidate feature to try again. This prevents infinite recursion / stack overflows in adversarial or flat objective landscapes. Larger values allow more exploration at a node (potentially better splits but slower), while smaller values make the builder more conservative (fewer retries, faster, more leaves). Default: 1000.

## Value

A list representing the tree/subtree.

---

<code>stratified_kfold</code>	<i>Stratified K-fold index generator</i>
-------------------------------	--

---

## Description

Splits indices into K folds while preserving the class distribution of a binary factor. This mimics scikit-learn's `StratifiedKFold`, ensuring each fold has a representative ratio of the two classes in S.

## Usage

```
stratified_kfold(S, K = 5)
```

## Arguments

<code>S</code>	A vector or factor indicating class membership (typically 0/1 or two-class factor) for stratification.
<code>K</code>	Integer number of folds ( $K \geq 2$ ). If $K$ is larger than the number of observations, it will be reduced to that number.

## Details

The function deterministically allocates indices to folds by class. For each class in S, indices are cyclically assigned to folds to balance counts. If  $K == 1$ , a single fold containing all indices is returned (though typically  $K$  should be  $\geq 2$  for cross-validation).

## Value

A list of length  $K$ , where each element is an integer vector of row indices assigned to that fold. The union of all folds equals `1:length(S)`, and folds are roughly equal in size.

---

<code>summary.characterizing_underrep</code>	<i>Summarize a characterizing_underrep fit</i>
--	--

---

## Description

Prints the ROOT summary (un/weighted estimates with standard deviations) and a brief overview of terminal rules from the annotated summary tree, if available.

## Usage

```
## S3 method for class 'characterizing_underrep'
summary(object, ...)
```

## Arguments

<code>object</code>	A <code>characterizing_underrep</code> object.
<code>...</code>	Unused; included for S3 compatibility.

## Details

Delegates core statistics to `summary(object$root)`; previews up to five terminal rules when a summary tree exists, and reports plot availability.

## Value

`object`, invisibly.

`summary.ROOT`

*Summarize a ROOT fit*

## Description

Summarizes a `ROOT` object by reporting the primary estimands and key model diagnostics. The first two lines report:

1. the **unweighted** estimate (ATE in RCT for single-sample or TATE for two-sample) and its standard deviation (SD);
2. the **weighted** estimate (Weighted ATE in RCT or WTATE, using the learned subgroup weights `w_opt`) and its SD.

Subsequent lines describe the estimand type, number of trees, size of the Rashomon set, presence of a summary tree, covariate count, observation count, baseline loss, selected-tree losses, and the proportion kept by `w_opt`.

## Usage

```
## S3 method for class 'ROOT'
summary(object, ...)
```

## Arguments

- |                     |  |
|---------------------|--|
| <code>object</code> | A <code>ROOT</code> object returned by <code>ROOT()</code> . |
| ...                 | Unused; included for S3 compatibility.                       |

## Details

The method prefers pre-computed estimates under `object$estimate` (as stored by `ROOT()`). If unavailable, it recomputes:

- Unweighted mean as  $\text{mean}(v)$  over the analysis set (all rows in single-sample;  $S=1$  in two-sample);
- Unweighted SD as  $\text{sd}(v)$  on the same set;
- Weighted mean as  $\sum wv / \sum w$  where  $w = w_{\text{opt}}$  (binary), and weighted SD as  $\sqrt{\sum w(v - \bar{v}_w)^2 / \sum w}$ .

The unweighted line corresponds to ATE in RCT/TATE, while the weighted line corresponds to Weighted ATE in RCT/WTATE over the selected subgroup.

## Value

The input `object`, *invisibly*. Printed output is a human-readable summary suitable for console use and logs.

---

<b>train</b>	<i>Train nuisance models for weighting</i>
--------------	--

---

## Description

Fits models to estimate sampling and treatment propensities on training data. Specifically, it computes:

- `pi`: The prevalence of sample inclusion (estimated as mean of  $S$ ).
- `pi_m`: A logistic regression model for  $P(S = 1 | X)$  using all covariates.
- `e_m`: A logistic regression model for  $P(Tr = 1 | X, S = 1)$ , fit only on the subset where  $S == 1$ .

## Usage

```
train(training_data, outcome, treatment, sample)
```

## Arguments

<code>training_data</code>	A data frame containing the training dataset.
<code>outcome</code>	Name of the outcome column (typically observed outcome, e.g. "Yobs").
<code>treatment</code>	Name of the treatment indicator column (e.g. "Tr").
<code>sample</code>	Name of the sample inclusion indicator column (e.g. "S").

## Details

This function uses simple logistic regression (`glm` with logit link) to estimate the necessary nuisance parameters for weighting. It requires that both  $S$  and  $Tr$  have variation (both 0 and 1 must be present in training data); if not, the fitting is not possible and an error is raised. All covariates other than the specified outcome, treatment, and sample columns are used as predictors.

## Value

A list with components:

<code>pi</code>	Numeric scalar giving the overall sample inclusion rate $P(S = 1)$ in the training data.
<code>pi_m</code>	A fitted <code>glm</code> model (binomial family) for $P(S = 1   X)$ .
<code>e_m</code>	A fitted <code>glm</code> model (binomial family) for $P(Tr = 1   X, S = 1)$ .

---

**train\_single***Train treatment propensity model (single-sample mode)*

---

**Description**

Fits a logistic regression for  $P(T = 1 | X)$  on the provided training data. Used by the single-sample Double ML path where no sample-selection model is needed.

**Usage**

```
train_single(training_data, outcome, treatment)
```

**Arguments**

- |                            |  |
|----------------------------|--|
| <code>training_data</code> | A data frame containing the outcome, treatment, and covariates. Only <code>treatment</code> and covariates are used for fitting. |
| <code>outcome</code>       | Name of the outcome column (character). Present for a consistent signature but not used in this function.                        |
| <code>treatment</code>     | Name of the binary treatment indicator column (0/1).   |

**Value**

A list with one element:

- |                  |  |
|------------------|--|
| <code>e_m</code> | A <code>glm</code> (binomial) object for the treatment propensity model. |
|------------------|--|

# Index

characterize\_tree, 2  
characterizing\_underrep, 3  
choose\_feature, 4  
  
estimate, 5  
estimate\_dml, 6, 6  
estimate\_dml\_single, 7  
estimate\_single, 7  
  
gen\_S, 8  
gen\_T, 8  
gen\_XY, 9  
get\_data, 10  
  
loss, 11  
loss\_from\_objective, 11  
  
midpoint, 12  
  
objective\_default, 11, 12  
objective\_if, 11, 13  
  
reduce\_weight, 13  
ROOT, 12, 14  
  
split\_node, 12, 15  
stratified\_kfold, 17  
summary.characterizing\_underrep, 17  
summary.ROOT, 18  
  
train, 6, 19  
train\_single, 7, 20